

Ceph RGW-Prefetching

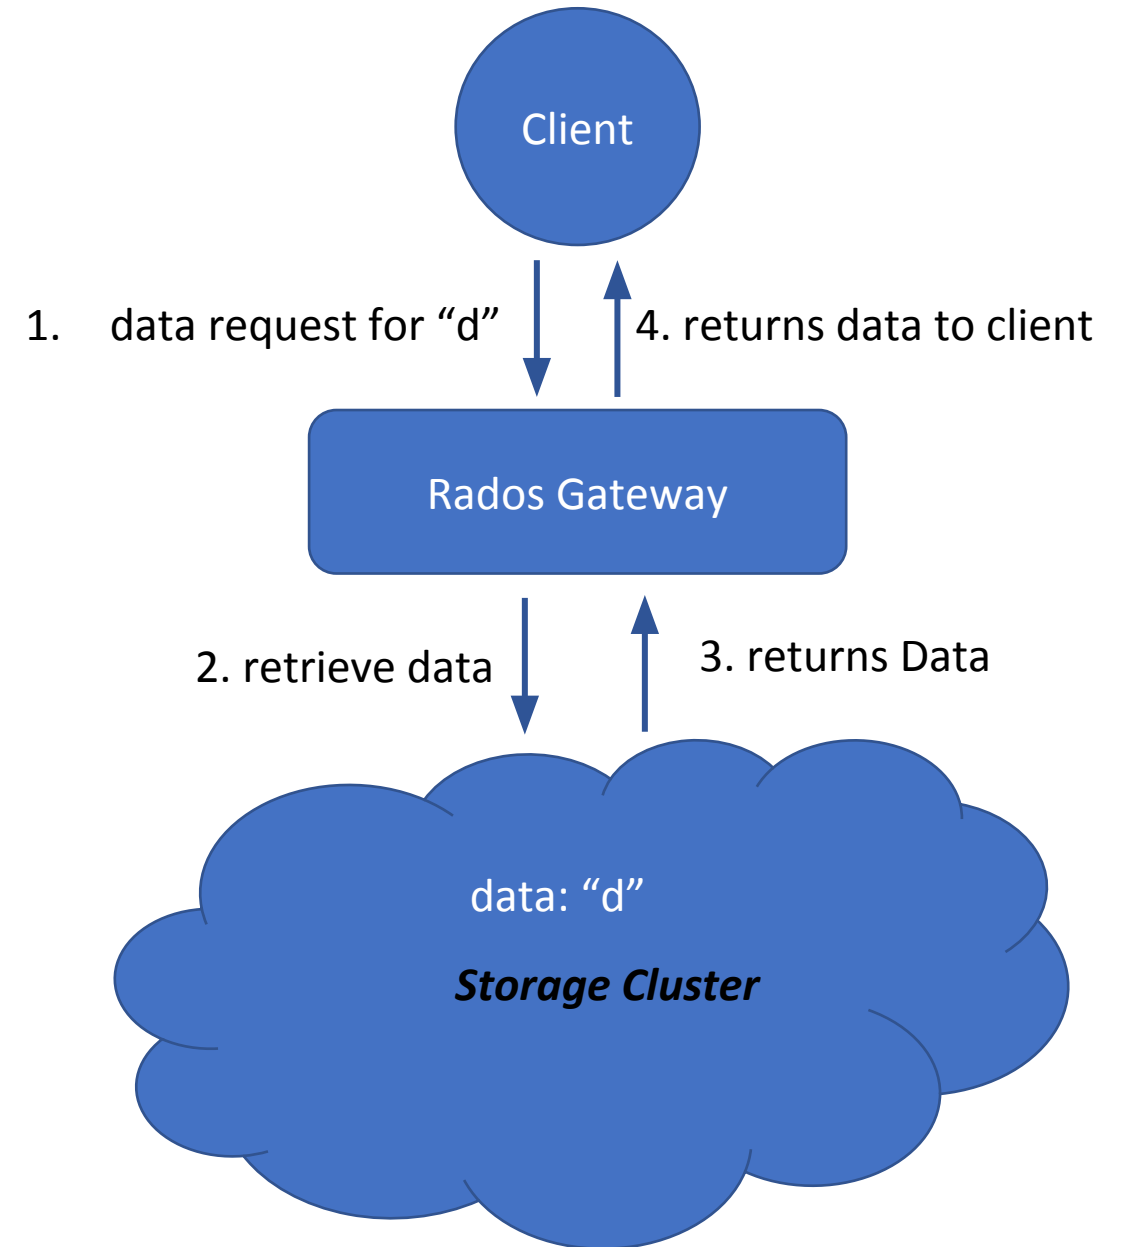
Ali R., Amin M., Bissenbay D.

What is CEPH?

- ❑ Distributed Storage System

- ❑ **Key Components** [Related to Project]

- ❑ Client (Applications)
- ❑ Rados Gateway
- ❑ Storage Cluster



Prefetching in Ceph

- ❑ Current caching scheme is re-active
- ❑ Goal: Improve performance in first access
 - ❑ Lower latency
 - ❑ Higher throughput
- ❑ Two type of prefetching
 - ❑ on-demand prefetching
 - ❑ pro-active prefetching in rgw

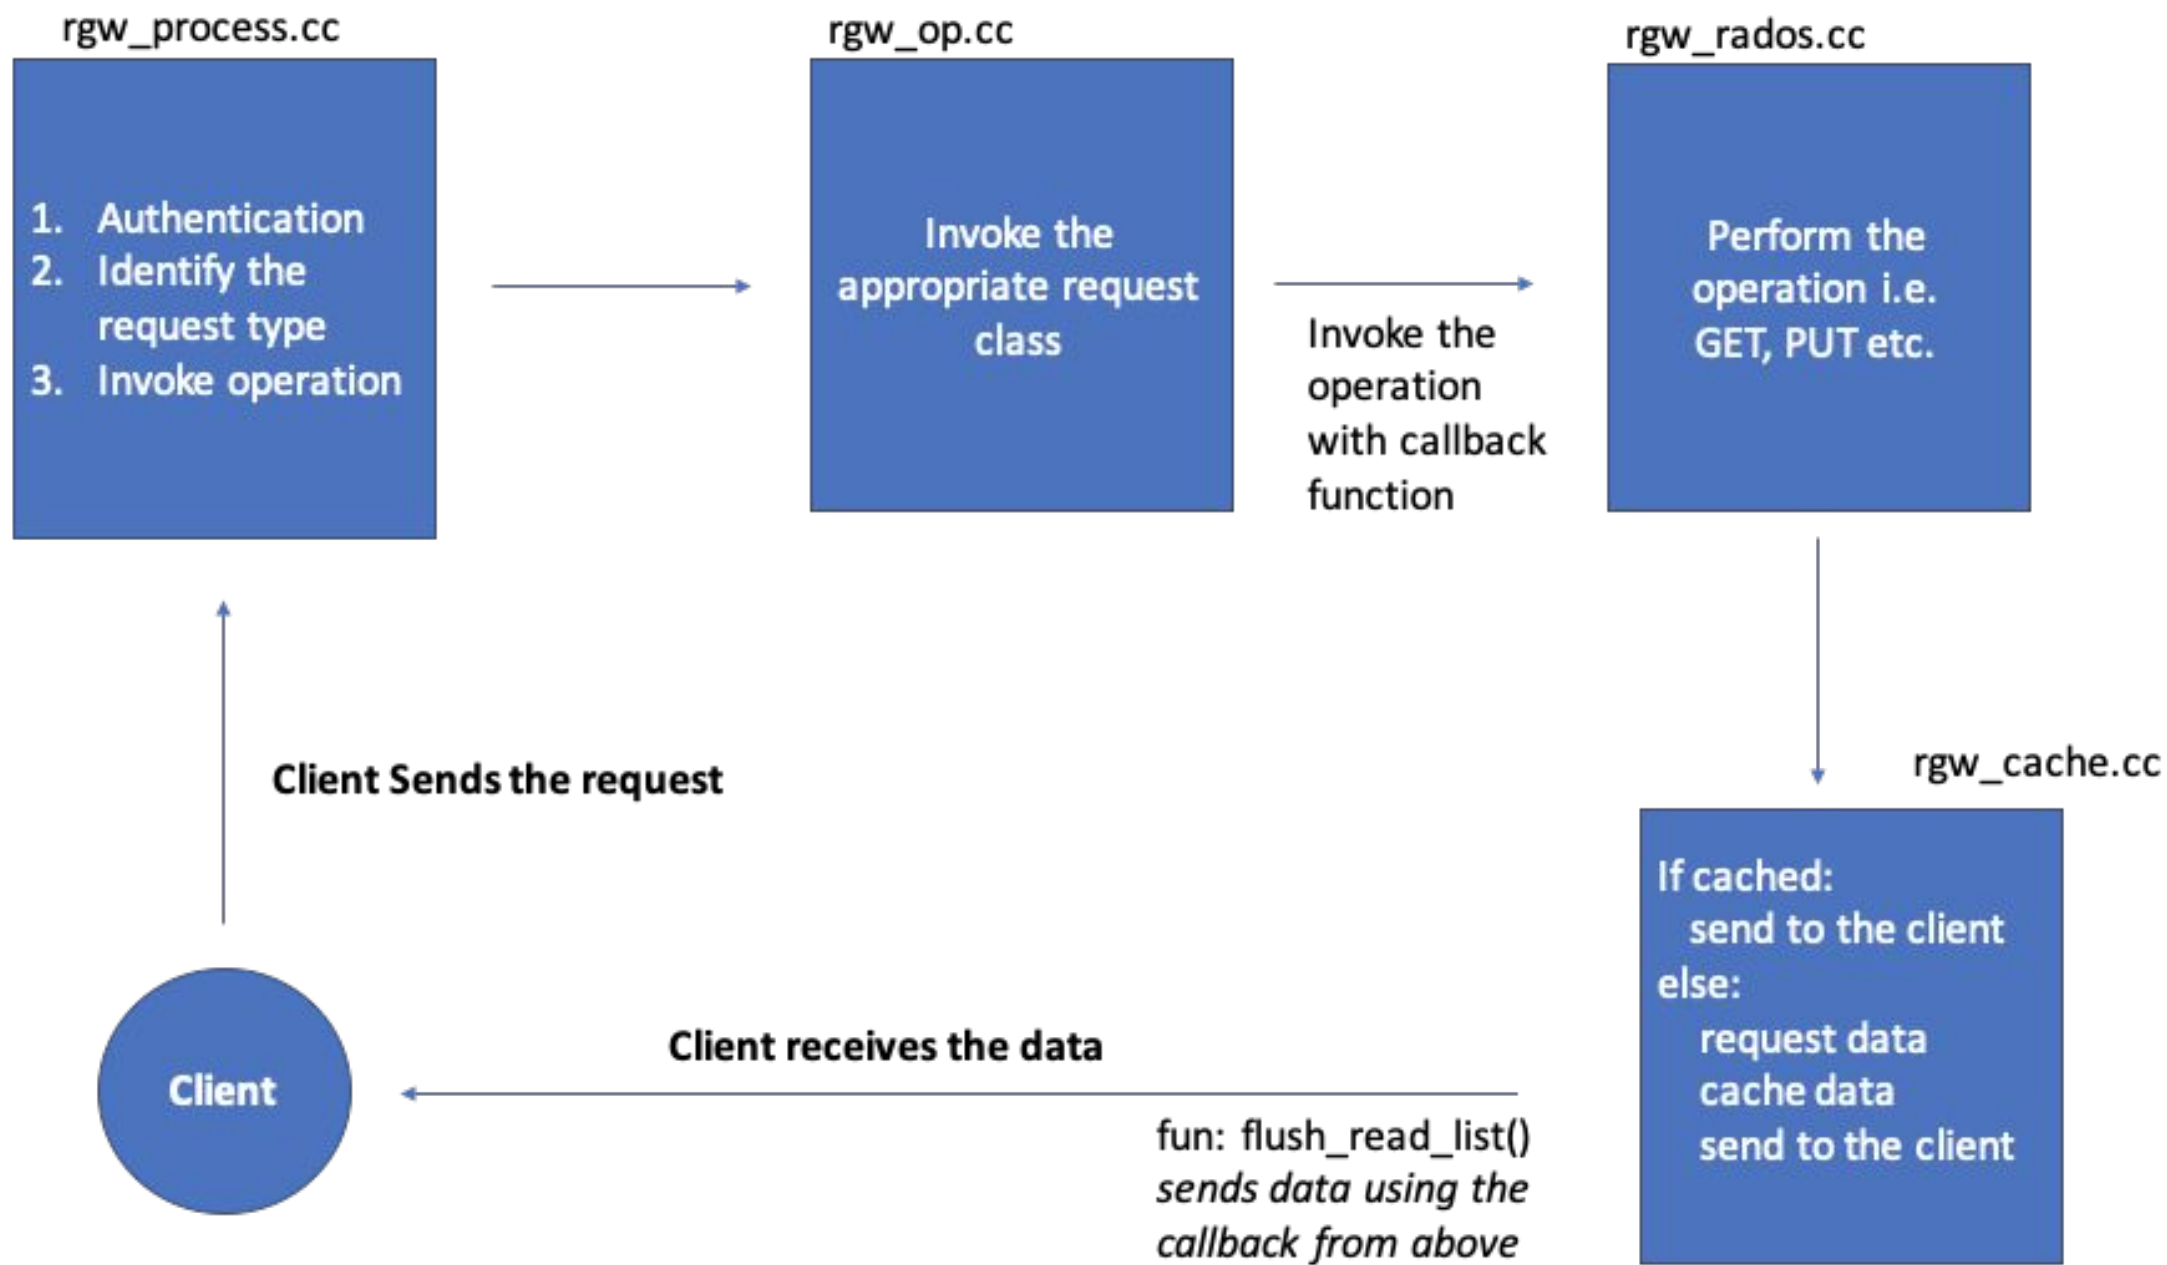
Until Sprint 3

- ❑ Uploaded & retrieved files to Ceph storage
 - ❑ S3cmd
 - ❑ Swift
 - ❑ Boto3
- ❑ Understanding the code
 - ❑ Deployed Ceph on multiple machines
 - ❑ Data flow
 - ❑ Debugging the code with GDB
 - ❑ Brainstorming with mentors to find the best development strategy
- ❑ In Progress:
 - ❑ Prefetching based on user request
 - ❑ Pro-active prefetching on block level
 - ❑ Evaluation

Sprint 4

- ❑ Done:
 - ❑ Prefetching based on user request
 - ❑ Pro-active prefetching on block level
- ❑ In progress:
 - ❑ Full testing of the above prefetching code
 - ❑ Evaluation (in terms of performance)
- ❑ Stretching the Goals:
 - ❑ Implement APIs to monitor the cache

Prefetching Based on User Request



Prefetching based on user request *[approach 1]*

- ❑ Overload the GET request (add header)
- ❑ Catch the header in rgw_rados.cc
 - ❑ Instead of sending back the data, place it in the cache
 - ❑ Reply the user with normal response headers (200 OK etc. and Content-length: 0)
- ❑ Issue with this approach
 - ❑ async data read requests get cancelled

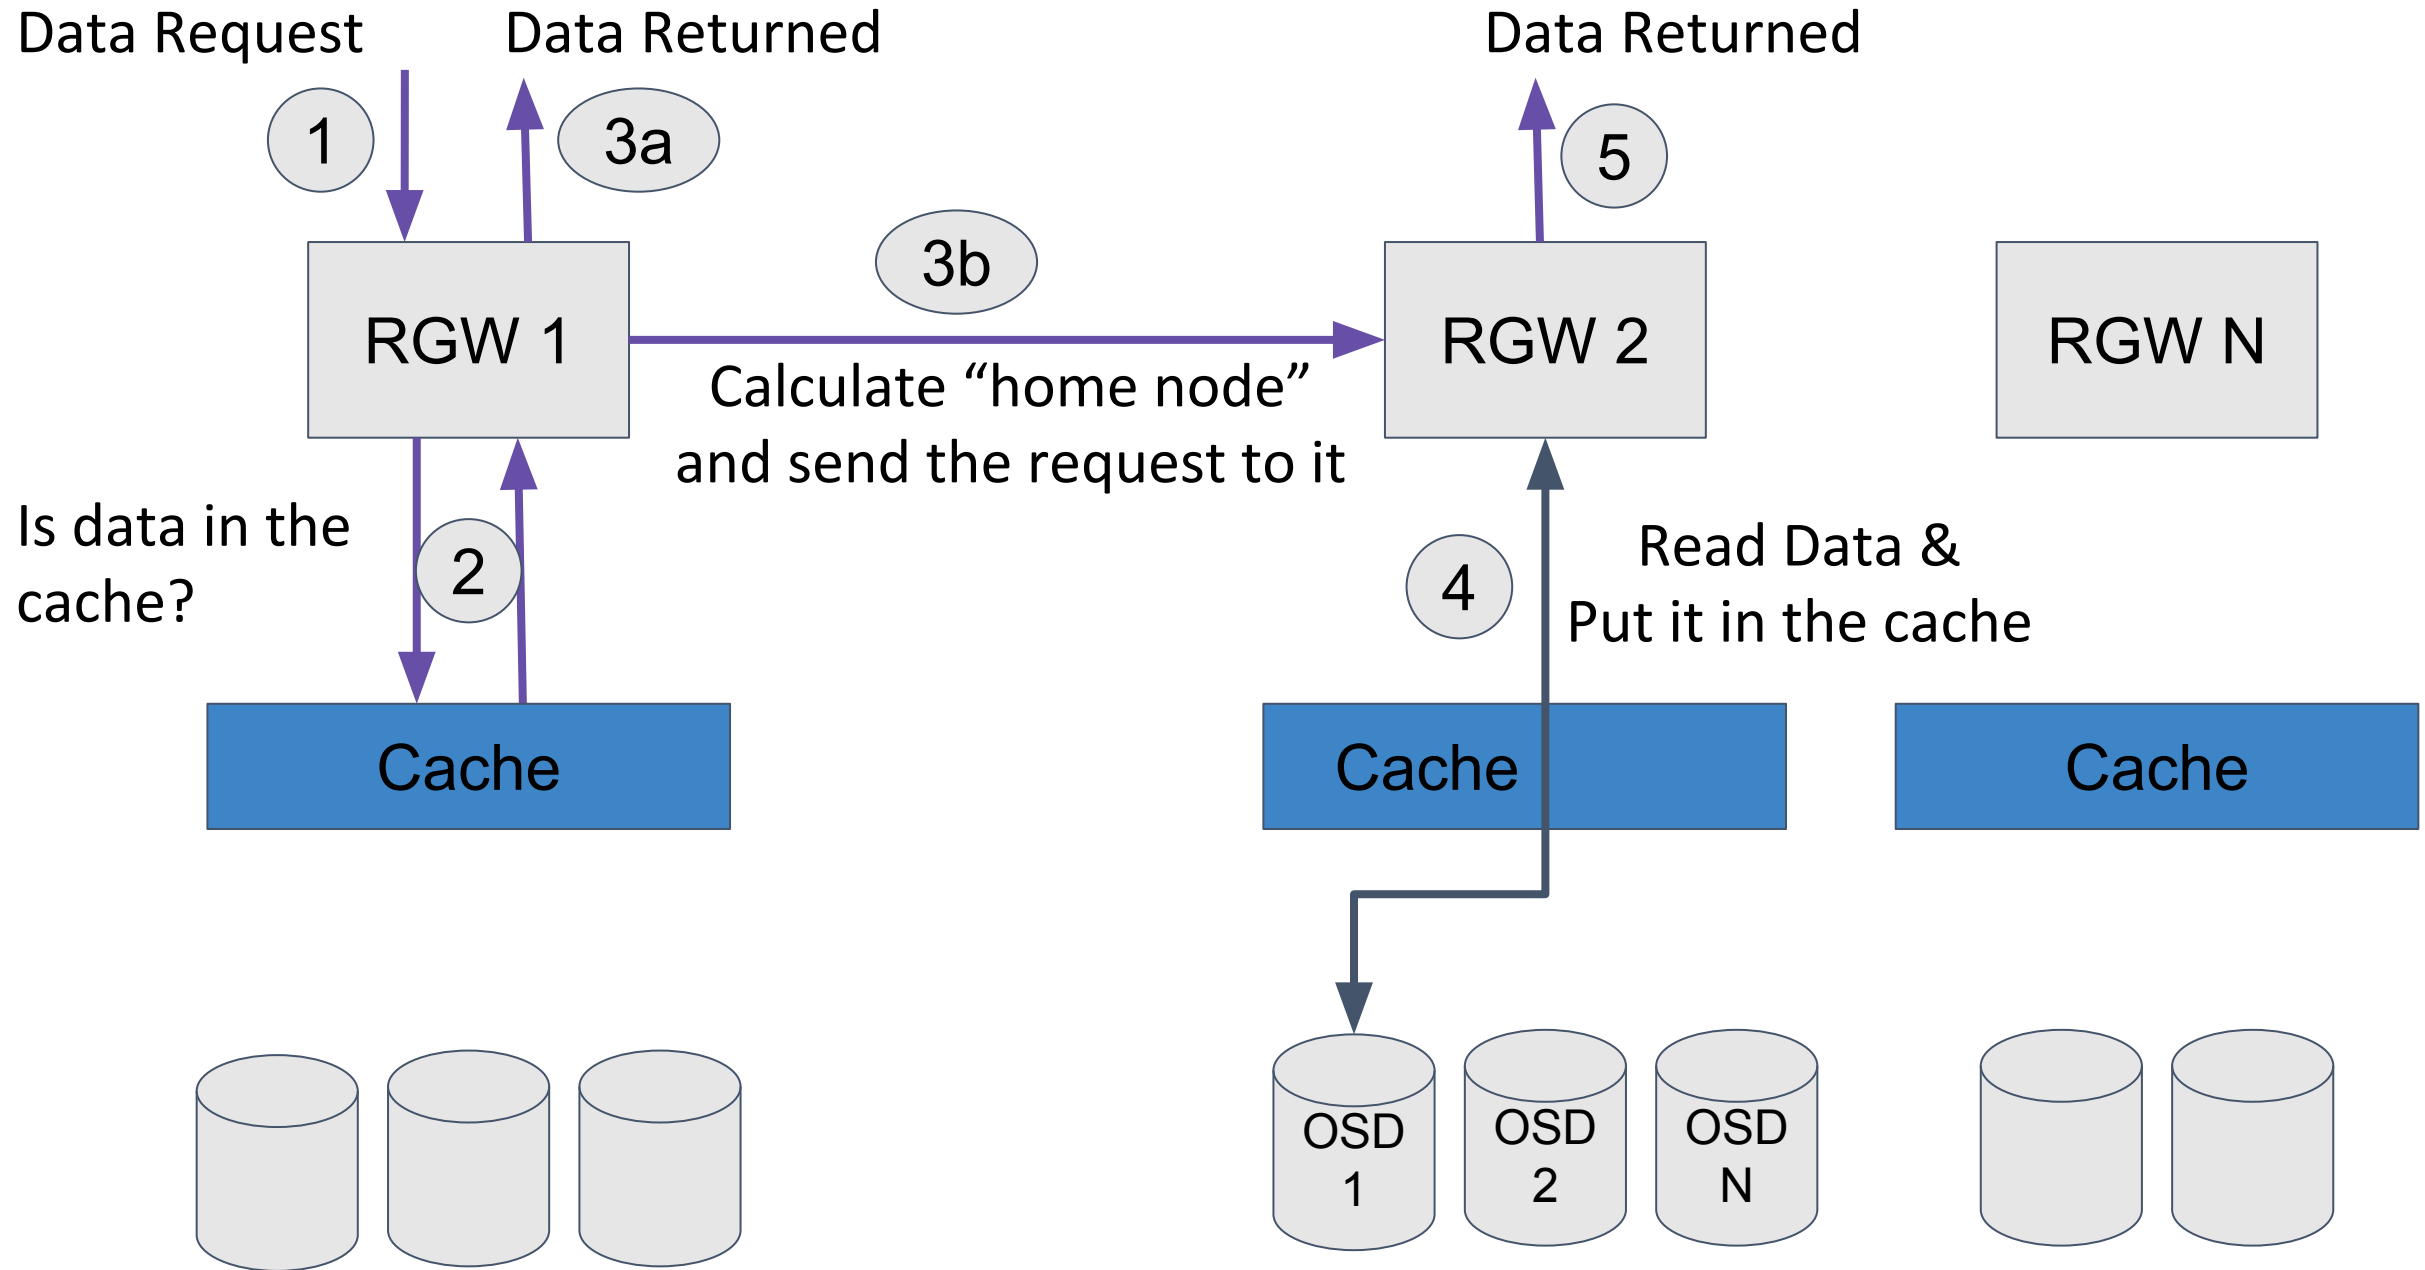
Prefetching based on user request *[approach 2]*

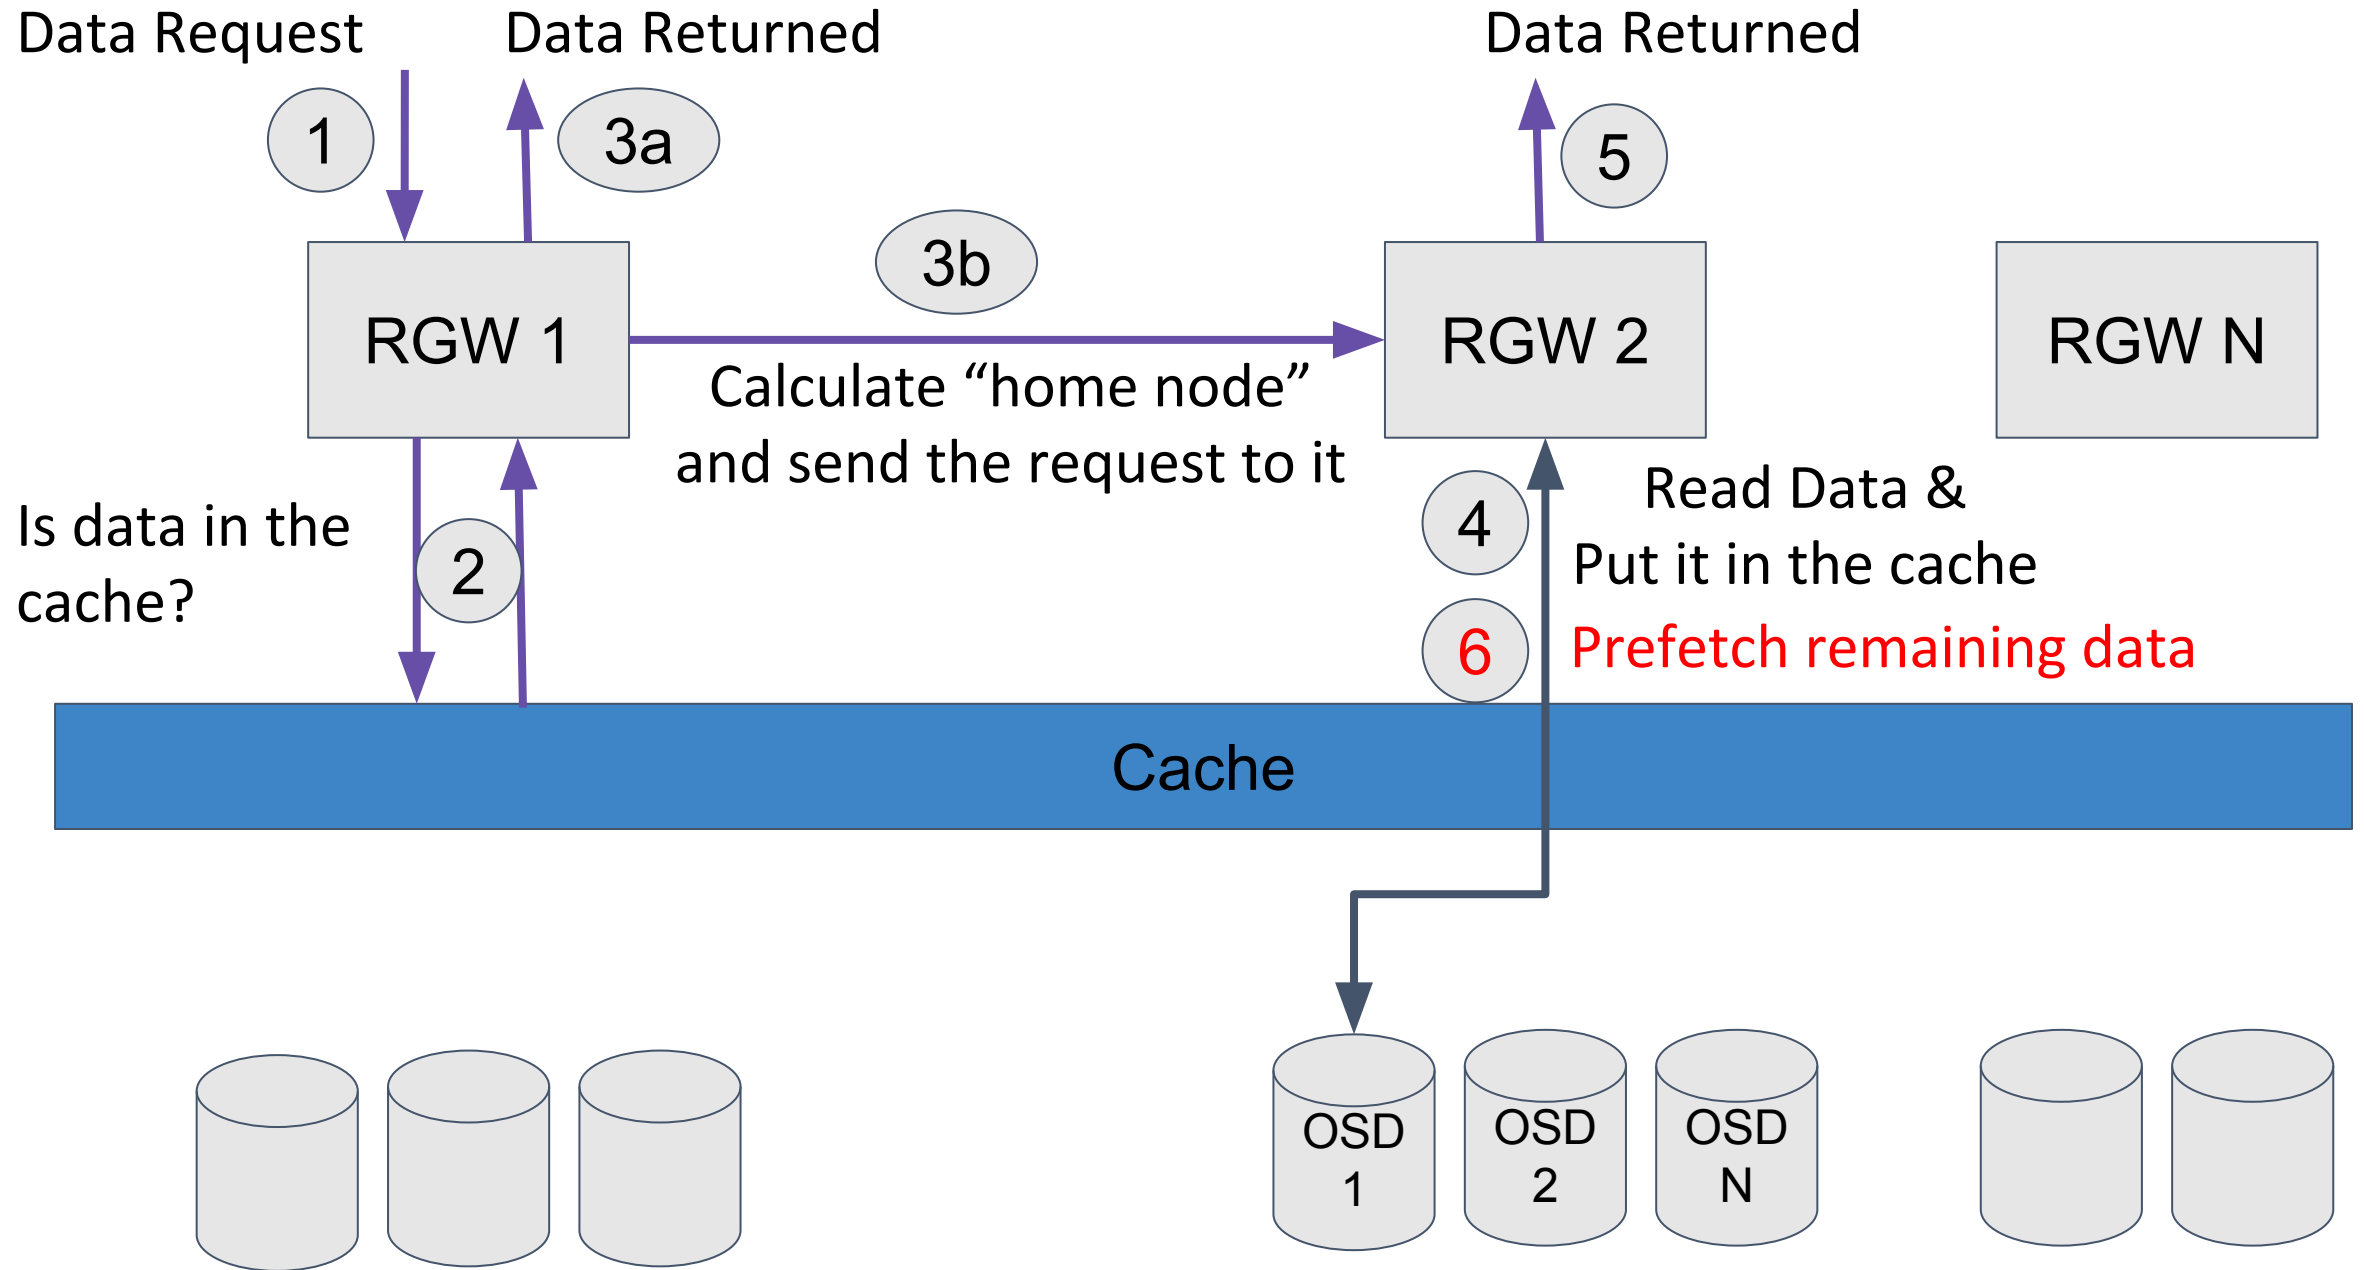
- ❑ Overload the GET request (add header)
- ❑ Catch the header in rgw_rados.cc
 - ❑ Instead of sending back the data, place it in the cache
 - ❑ **Wait for all of the data to be written in the cache and then**
 - ❑ Reply the user with normal response headers (200 OK etc. and Content-length: **x bytes**)
 - ❑
- ❑ Issue with this approach
 - ❑ It's a blocking call

What's next for “on-demand prefetching”

- ❑ Talked to our mentors
- ❑ Approach 1 is more logical
- ❑ We will find a way that data-request are not cancelled even though client closes the connection

Pro-active Prefetching





Pro-active prefetching

- ❑ What do we need?
 - ❑ What data?
 - ❑ Where is it?
 - ❑ What is the size of it?
 - ❑ What chunk has the user read?
 - ❑ We should not block other tasks
 - ❑ We should not harm the performance

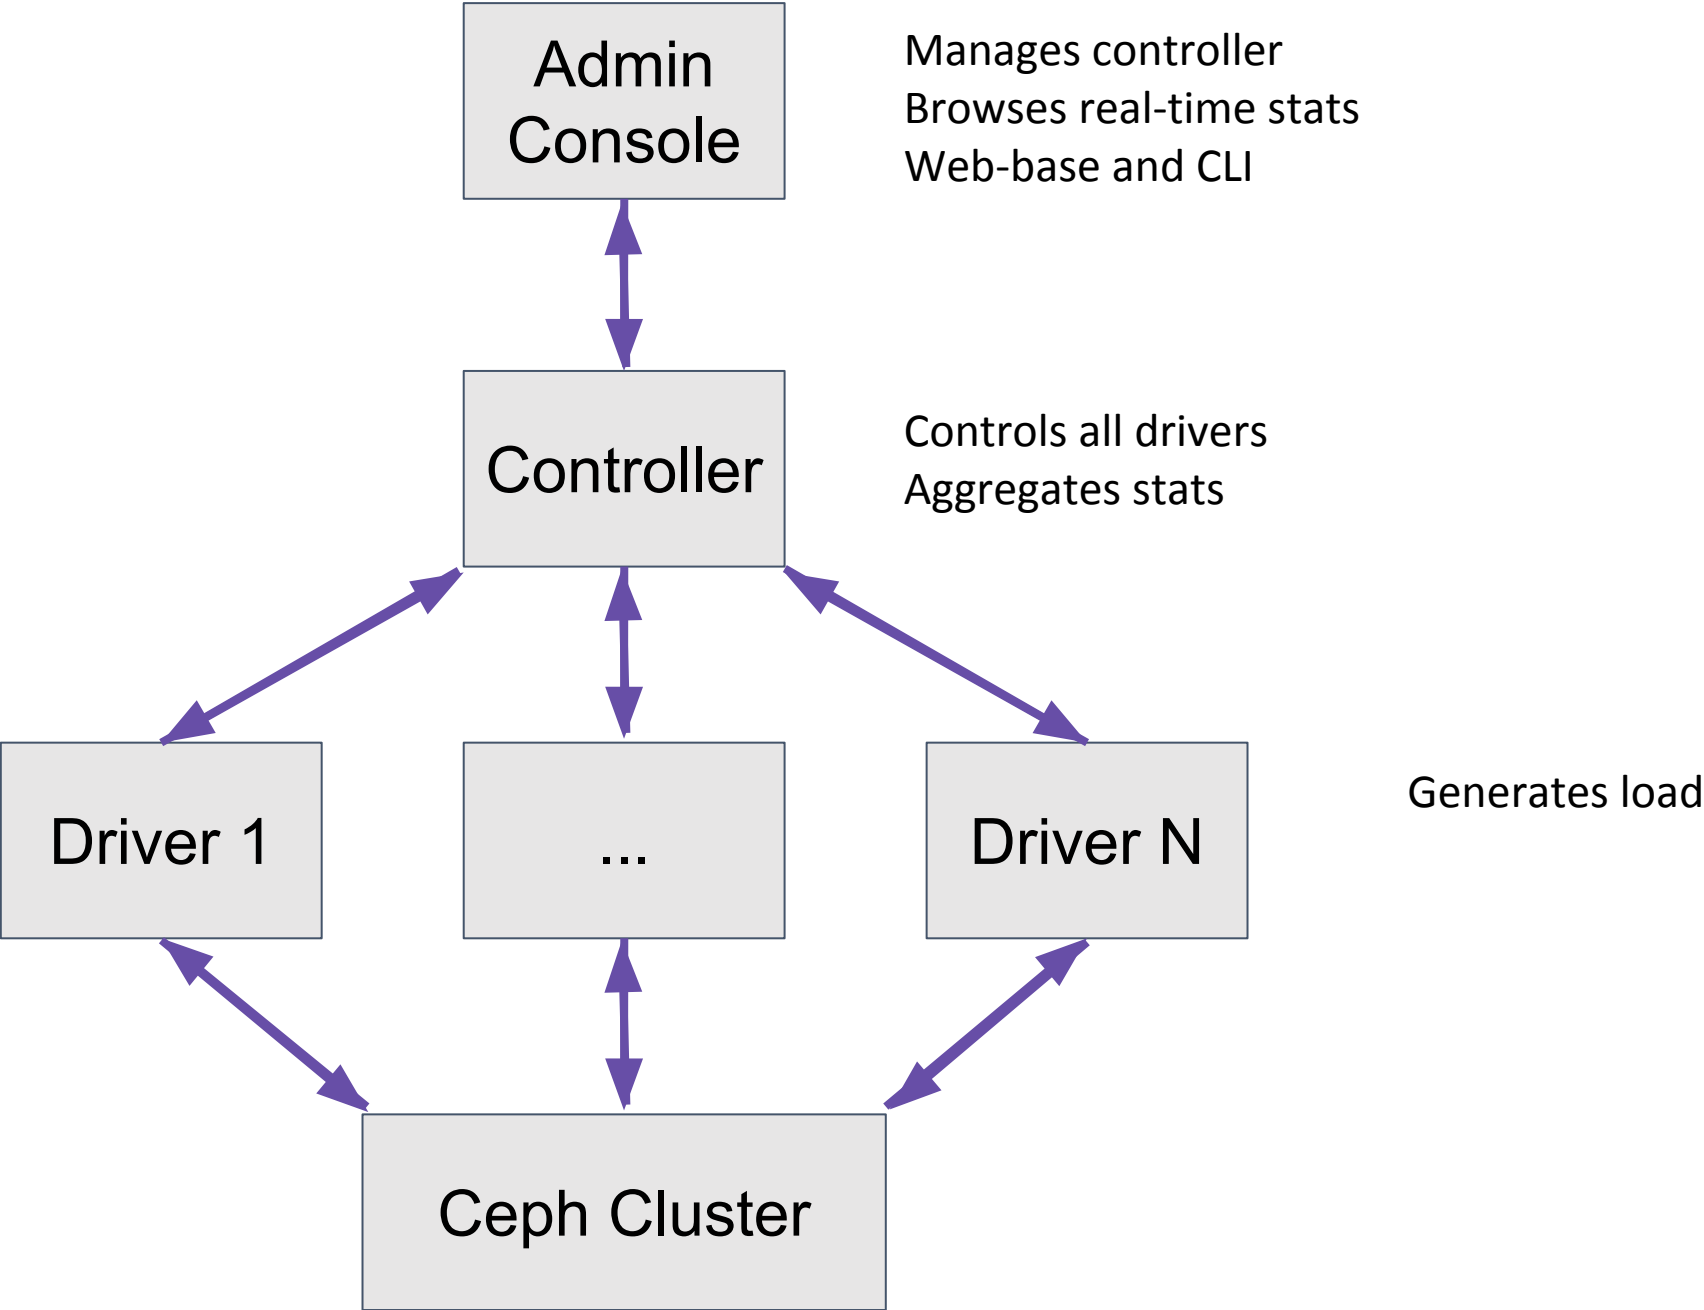
Pro-active prefetching

- ❑ A prefetching C++ class to talk with cache
 - ❑ We have implemented:
 - ❑ The requested File, size of requested chunk and its offsets
 - ❑ The size of the File and passing it to the prefetching class
 - ❑ User's authentication
- ❑ A thread pool for prefetching
- ❑ Adding prefetching task to the thread pool

Evaluation

- ❏ COSBench
- ❏ Modifying COSBench

COSBench



Workload Configuration

```
<?xml version="1.0" encoding="UTF-8" ?>
<workload name="s3-sample" description="sample benchmark for s3">

  <storage type="s3" config="accesskey=;secretkey=;endpoint=http://localhost:8000" />

  <workflow>
    <workstage name="init">
      <work type="init" workers="1" config="cprefix=s3testqwer;containers=r(1,2)" />
    </workstage>

    <workstage name="prepare">
      <work type="prepare" workers="1" config="cprefix=s3testqwer;containers=r(1,2);objects=r(1,10);sizes=c(64)KB" />
    </workstage>

    <workstage name="main">
      <work name="main" workers="1" runtime="30">
        <operation type="read" ratio="50" config="cprefix=s3testqwer;containers=u(1,2);objects=u(1,10)" />
        <operation type="write" ratio="50" config="cprefix=s3testqwer;containers=u(1,2);objects=u(11,20);sizes=c(64)KB" />
      </work>
    </workstage>

    <workstage name="cleanup">
      <work type="cleanup" workers="1" config="cprefix=s3testqwer;containers=r(1,2);objects=r(1,20)" />
    </workstage>

    <workstage name="dispose">
      <work type="dispose" workers="1" config="cprefix=s3testqwer;containers=r(1,2)" />
    </workstage>
  </workflow>
</workload>
```

Flexible load control

Object size distribution

Read/Write Operations

```
<operation type="prefetch" ratio="50" config="cprefix=s3testqwer;containers=u(1,2);objects=u(1,10)" />
```

Evaluation

❑ Done:

- ❑ Installed COSBench on the machine
- ❑ Run a testing workload
- ❑ Agreed on design change of a xml config file with mentors

❑ In progress:

- ❑ Setting Up Dev Environment in Eclipse to make changes in COSBench
- ❑ Creating a Prefetching Operator which will send headers on GET

❑ Next Step:

- ❑ Come up with a workload design that will be used to test prefetching
- ❑ Try different workloads to identify bottlenecks

Burndown chart

